

1. JOINS: Used to Collect Data from Two or More Tables

- Cartesian Product: when there is no relation between two tables:
Each row in the first table is paired with all rows in the second table
Formed when the join condition is omitted
To avoid this, always include a valid JOIN condition

```
SELECT
    last_name,
    name AS department_name
FROM
    employees,
    departments;
```

- Inner Join: produces a set of records that match in both tables

```
SELECT
    e.first_name,
    d.name
FROM employees AS e
INNER JOIN departments --or just JOIN
ON e.department_id = d.department_id
```

- Left Join: matches every entry in left table regardless of match in the right

```
SELECT
    e.first_name,
    d.name
FROM employees AS e
LEFT JOIN departments
ON e.department_id = d.department_id
```

- Right Join: matches every entry in right table regardless of match in the left

```
SELECT
    e.first_name,
    d.name
FROM employees AS e
LEFT JOIN departments
ON e.department_id = d.department_id
```

- Full Join: returns all records in both tables regardless of any match

```
SELECT
    e.first_name,
    d.name
FROM employees AS e
FULL JOIN departments
ON e.department_id = d.department_id
```

- Cross Join: produces a set of associated rows of two tables
Multiplication of each row in the first table with each row in the

second one

The result is a cartesian product where there is no condition in WHERE clause

```
SELECT
    d.department_id,
    d.name,
    e.employee_id,
    e.first_name
FROM
    departments AS d
CROSS JOIN employees as e;
```

2. Subqueries: query manipulation on multiple level

- Subqueries: SQL query inside a larger one
- Can be nested in SELECT, INSERT, UPDATE, DELETE
- Usually added within a WHERE clause

```
SELECT
    employee_id AS id,
    first_name,
    department_ud
FROM
    employees
WHERE
    department_id = 1;
```

```
SELECT
    COUNT(e.employee_id) AS count
FROM
    employees as e
WHERE e.salary >
(
    SELECT
        AVG(salary) AS "average_salary"
    FROM
        employees
);
```

3. Indices: Speed Retrieval of Rows

- Special lookup tables that speed retrieval of rows
- Usually implemented as B-trees
- Speed up SELECT queries and WHERE queries
- Slows down data input
- Should be used for big tables only(e.g.: 50 000rows)
- Should not be used on columns that contain a high number of NULL values

```
CREATE INDEX index_name_index
ON table_name(first_column, second_column);
```

****clustered and non-clustered****